

Fuzzing at Mach Speed

Uncovering MacOS and iOS IPC Vulnerabilities with Dillon Franke



Who Am I?



CURRENTLY

Senior Proactive Security Consultant
(Pentesting)

*Application Security
Source Code Reviews
Embedded Device Assessments*

Project Zero 20% Project

PREVIOUSLY

FLARE Offensive Task Force (OTF)
(Reverse Engineering)

*Malware reversing
Searching for exploits used in the wild
0-day vulnerability research
Exploit development*

STUDIED


Bachelor's & Master's in Computer Science at Stanford University
Security and Systems Engineering

HOBBIES

Playing Guitar
Cycling in the San Francisco Bay Area
Hacking (obviously)



- Offensive Security Researchers
- Defensive Security Engineers
- Software Developers
- Mobile Application Researchers



**Who is this
talk for?**

Overview

*Join me as I dive into my process
searching for low-level vulnerabilities
in MacOS over the past year.*



Crash Course on Fuzzing and IPC Mechanisms



The Fuzzing Cycle



Next Steps




Q&A



CRASH COURSE

What is Fuzzing?



Fuzzing is sending unexpected
inputs to a **system** in the hopes of
making something unexpected
happen





CRASH COURSE


What is an Attack Vector?



An attack vector is a channel to send an **input** to a **system**



Bluetooth



Interprocess
Communications



Notifications



Peripherals




Wireless Connection





CRASH COURSE

What is an Attack Vector?



An attack vector is a channel to send an **input** to a **system**

Adobe Acrobat


Open PDF Functionality

Google Search

Query Parameter
(<https://google.com?query=<INPUT>>)

Smart Watch

Bluetooth Data Handling





CRASH COURSE Why Fuzz?

- 1 In memory-unsafe languages, (C/C++) we want to send input that causes a crash

- 2 Depending on the type of crash, our input might be able to trigger:
 - Buffer Overflow
 - Heap Overflow
 - Use-After-Free
 - Double Free
 - Memory Leak (bypass ASLR)

November 15th, 2023

Adobe Acrobat Reader DC Font Parsing Use-After-Free Remote Code Execution Vulnerability

ZDI-23-1690

ZDI-CAN-21929

CVE ID CVE-2023-44367

CVSS SCORE 7.8, (AV:L/AC:L/PR:N/UI:R/S:U/C:H/I:H/A:H)

AFFECTED VENDORS Adobe

AFFECTED PRODUCTS Acrobat Reader DC

VULNERABILITY DETAILS This vulnerability allows remote attackers to execute arbitrary code on affected installations of Adobe Acrobat Reader DC. User interaction is required to exploit this vulnerability in that the target must visit a malicious page or open a malicious file.

Use-After-Free

2

The specific flaw exists within the parsing of embedded fonts. The issue results from the lack of validating the existence of an object prior to performing operations on the object. An attacker can leverage this vulnerability to execute code in the context of the current process.

1 Attack Vector





CRASH COURSE

Different Types of Fuzzing

Mutation-Based

Fuzzing: Modify existing inputs to create new ones, then send them to the program


Grammar-Based

Fuzzing: Generate inputs based on specified rules defining the structure of valid inputs



CRASH COURSE

What is the XNU Kernel?

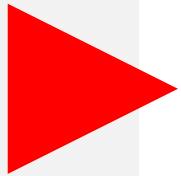


XNU (X is Not Unix) is the kernel that powers macOS.



CRASH COURSE

What is the XNU Kernel?



XNU (X is Not Unix) is the kernel that powers macOS.

Mach Layer: Responsible for low-level tasks like thread management, interprocess communication (IPC), and memory management.


BSD Layer: Handles higher-level POSIX tasks, like file system, network, and security.

I/O Kit: A framework for developing device drivers, designed with a model resembling object-oriented programming.



CRASH COURSE

What is the XNU Kernel?



XNU (X is Not Unix) is the kernel that powers macOS.

Mach Layer: Responsible for low-level tasks like thread management, **interprocess communication (IPC)**, and memory management.

BSD Layer: Handles higher-level POSIX tasks, like file system, network, and security.


I/O Kit: A framework for developing device drivers, designed with a model resembling object-oriented programming.




CRASH COURSE

What is Interprocess Communication (IPC)?

Processes need to talk to each other!



How do they do this?



Mach Messages



- Mach exceptions
- Unix signals
- Unnamed pipes
- Named pipes (fifos)
- XSI/System V IPC
- POSIX IPC
- Distributed Objects
- Apple Events
- Core Foundation IPC mechanisms



CRASH COURSE

What is Interprocess Communication (IPC)?

Processes need to talk to each other!



How do they do this?



Mach Messages


- Mach exceptions
- Unix signals
- Unnamed pipes
- Named pipes (fifos)
- XSI/System V IPC
- POSIX IPC
- Distributed Objects
- Apple Events
- Core Foundation IPC mechanisms

Lowest level IPC mechanism and
the direct basis for many higher
level mechanisms



CRASH COURSE

What are Mach Ports?



IPC message queues, managed by the kernel

Port Right: Handle to a port that allows sending or receiving messages to the port

Receive Right: Allows receiving a mach port's messages

Send Right: Allows sending messages to a mach port

Send Once: Allows sending a single message to a mach port



CRASH COURSE

What are Mach Ports?

IPC message queues,
managed by the
kernel

Port Right: Handle to a port that allows sending or receiving messages to the port

Receive Right: Allows receiving a mach port's messages

Send Right: Allows sending messages to a mach port

Send Once: Allows sending a single message to a mach port

> lsmp -h

Usage: lsmp -p <pid> [-a|-v|-h]

Lists information about mach ports.
Please see man page for description of each column.

Process (135) : kextd													
name	ipc-object	rights	flags	boost	reqs	recv	send	sonce	oref	qlimit	msgcount	context	
0x000000103	0xdce4a79b	send	-----	---	---	1	2			5	0	0x0000000000	
0x000000203	0xdd0c45e3	recv	-----	0	---	1		54		N	5	0	0x0000000000
0x000000307	0xd6247d5b	send	-----	---	0	1				N	5	0	0x0000000000
0x000000403	0xdd0c41f3	recv	-----	---	0	1				N	5	0	0x0000000000
0x000000503	0xdd0c564b	recv	-----	---	0	1				N	5	0	0x0000000000
0x000000603	0xdce4a8eb	send	-----	---	---	1				N	5	0	0x0000000000
0x000000703	0xdd0c56f3	recv	-----	0	---	1				N	5	0	0x0000000000
0x000000803	0xd624781b	send	-----	---	---	1				N	5	0	0x0000000000
0x000000903	0xdcc335a3	recv, send	--GS--	0	---	1	2			Y	5	0	0x0000000000
0x000000a03	0xdcc690e3	recv, send	--GS--	0	---	1	1			Y	5	1	0x0000000000
	+ send		-----	---	---	1				<-			
0x000000b03	0xdcc6957b	send	-----	---	---	1				->	1	0	0x0000000000
0x000000c03	0xdcc69623	send	-----	---	---	1				->	1	0	0x0000000000
0x000000d0f	0xde2da7db	recv	-----	0	---	1				Y	5	0	0x0000000000
	+ send		-----	D--	---	1				<-			
0x000000e07	0xd6248fb	send	-----	---	---	1				->	32	0	0x0000000000
0x000000f03	0xdcaeef13	send	-----	---	---	1				->	6	0	0x0000000000
0x000001003	0xdcaefbcb	send	-----	---	---	1							
0x000001103	0xd6247e03	send	-----	---	---	1							
0x000001203	0xdcc6abcb	recv, send	-----	0	---	1	1			Y	5	0	0x0000000000
0x000001303	0xd779214b	send	-----	---	---	6				->	128	0	0x0000000000
0x000001403	0xdd0c2cf3	send	-----	---	---	1							
0x000001507	0xdcbe5718b	send	-----	---	---	1				->	6	0	0x0000000000

```
total      = 845
SEND      = 841
RECEIVE   = 5
SEND_ONCE  = 0
PORT_SET   = 0
DEAD_NAME  = 0
DNREQUEST = 0
VOUCHERS  = 0
```

Single Process!!



Establishing a Mach Connection

Bootstrap Server

- A mach port to help establish connections with other mach ports
- By default, all processes have a send right to the bootstrap server

Mach Service

- A mach port with a name that is registered with the Bootstrap Server (e.g. **com.apple.troopers**)

Communicating with a Service

- 1 Alice allocates a new mach port, giving her a receive right
- 2 Alice registers her service using a specific name **com.apple.troopers** with the bootstrap server
By registering, Alice is giving the bootstrap server a send right to the port Alice has a receive right to
- 3 Bob asks the bootstrap server for the service named **com.apple.troopers** and the server gives Bob a copy of the send right for Alice's mach port
- 4 Bob can now send messages to Alice's mach port for Alice to receive



What are Mach Messages?

A struct used to exchange data between mach ports

Sending/Receiving Mach Messages

```
/*
 * Routine: mach_msg
 * Purpose:
 * Send and/or receive a message. If the message operation
 * is interrupted, and the user did not request an indication
 * of that fact, then restart the appropriate parts of the
 * operation silently (trap version does not restart).
*/
__WATCHOS_PROHIBITED __TVOS_PROHIBITED
extern mach_msg_return_t mach_msg(
    mach_msg_header_t *msg,
    mach_msg_option_t option,          Option specifies
    mach_msg_size_t send_size,         send/receive!
    mach_msg_size_t recv_size,
    mach_port_name_t recv_name,
    mach_msg_timeout_t timeout,
    mach_port_name_t notify
);
```




CRASH COURSE


What are Mach Messages?

A struct used to exchange data between mach ports

Sending/Receiving Mach Messages

```
/*
 * Routine: mach_msg
 * Purpose:
 * Send and/or receive a message. If the message operation
 * is interrupted, and the user did not request an indication
 * of that fact, then restart the appropriate parts of the
 * operation silently (trap version does not restart).
 */
__WATCHOS_PROHIBITED __TVOS_PROHIBITED
extern mach_msg_return_t mach_msg(
    mach_msg_header_t *msg,
    mach_msg_option_t option,           Option specifies
    mach_msg_size_t send_size,          send/receive!
    mach_msg_size_t recv_size,
    mach_port_name_t recv_name,
    mach_msg_timeout_t timeout,
    mach_port_name_t notify
);
```





The Fuzzing Cycle


Identify an
attack vector

Generate a
Corpus of
Inputs

Create a
Fuzzing
Harness

Fuzz and
Produce
Crashes

Identify
Relevant
Crashes



The Fuzzing Cycle

Abusing Mach Messages

Identify an attack vector


Generate a Corpus of Inputs

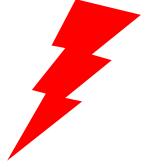
Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Sandbox Escape





The Fuzzing Cycle

Abusing Mach Messages

Identify an attack vector


Generate a Corpus of Inputs


Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Sandbox Escape





The Fuzzing Cycle

Abusing Mach Messages

Identify an attack vector


Generate a Corpus of Inputs


Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Sandbox Escape





The Fuzzing Cycle

Abusing Mach Messages

Identify an attack vector


Generate a Corpus of Inputs

Create a Fuzzing Harness


Fuzz and Produce Crashes

Identify Relevant Crashes

Sandbox Escape



Privilege Escalation





The Fuzzing Cycle

Abusing Mach Messages

Identify an attack vector


Generate a Corpus of Inputs

Create a Fuzzing Harness


Fuzz and Produce Crashes


Identify Relevant Crashes

Sandbox Escape



Privilege Escalation





The Fuzzing Cycle

Abusing Mach Messages

Identify an attack vector


Generate a Corpus of Inputs

Create a Fuzzing Harness


Fuzz and Produce Crashes


Identify Relevant Crashes

Sandbox Escape



Privilege Escalation





The Fuzzing Cycle

Previous Mach Research

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

BlackHat: Breaking the Chrome Sandbox with Mojo

- <https://i.blackhat.com/USA-22/Wednesday/US-22-Roettger-Breaking-the-Chrome-Sandbox-with-Mojo.pdf>
- Race condition + DoS == RCE

A Methodical Approach to Browser Exploitation

- <http://blog.ret2.io/2018/06/05/pwn2own-2018-exploit-development/>
- Safari sandbox escape via mach IPC messages == RCE



The Fuzzing Cycle

Finding Sandbox-Allowed Communications

How do we know what processes could allow an escape?

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

sbtool: <https://newosxbook.com/src.jl?tree=listings&file=/sbtool.c>

- Use built-in **sandbox_check()** function to determine which mach services a process can send to
- Message handlers we can send to → potential for sandbox escapes

```
> ./sbtool 2813 mach
com.apple.logd
com.apple.xpc.smd
com.apple.remoted
com.apple.metadata.mds
com.apple.coreduetd
com.apple.apsd
com.apple.coreservices.launchservicesd
com.apple.bsd.dirhelper
com.apple.logind
com.apple.revision
...Truncated...
```



The Fuzzing Cycle

Finding Sandbox-Allowed Communications

How do we know what processes could allow an escape?

Identify an attack vector

Generate a Corpus of Inputs


Create a Fuzzing Harness


Fuzz and Produce Crashes

Identify Relevant Crashes

sbtool: <https://newosxbook.com/src.jl?tree=listings&file=/sbtool.c>

- Use built-in **sandbox_check()** function to determine which mach services a process can send to
- Message handlers we can send to → potential for sandbox escapes





The Fuzzing Cycle

Finding an Entry Point

Identify an attack vector

Generate a Corpus of Inputs


Create a Fuzzing Harness


Fuzz and Produce Crashes

Identify Relevant Crashes

1 We know that **mach_msg()** is used to send mach messages from one process to another

2 Why not just modify real mach messages being sent?





The Fuzzing Cycle

Finding an Entry Point

Identify an attack vector

Generate a Corpus of Inputs


Create a Fuzzing Harness


Fuzz and Produce Crashes

Identify Relevant Crashes

1 We know that **mach_msg()** is used to send mach messages from one process to another

2 Why not just modify real mach messages being sent?





The Fuzzing Cycle

Finding an Entry Point

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes


Identify Relevant Crashes


Pros:

- Simple
- Similar to end exploit

Cons:

- Slow (*At mercy of the application to send messages*)
- Many points of potential failure
- Two different process spaces (code coverage difficult)
- Difficult to determine which message caused crash





The Fuzzing Cycle

Finding an Entry Point

Identify an attack vector


Generate a Corpus of Inputs

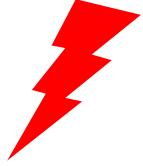
Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Instead of waiting for **mach_msg()** to be called,
what if we write a program to call it ourselves?





The Fuzzing Cycle

Finding an Entry Point

Identify an attack vector


Generate a Corpus of Inputs

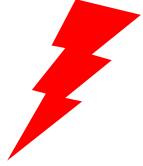
Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Instead of waiting for **mach_msg()** to be called, what if we write a program to call it ourselves?





The Fuzzing Cycle

Finding an Entry Point

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes


Identify Relevant Crashes


Instead of waiting for **mach_msg()** to be called, what if we write a program to call it ourselves?

Even Better: What if we just called the message handler directly?

*Some reverse engineering required

Getting "close" to the system of interest





The Fuzzing Cycle

Finding an Entry Point

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes


Identify Relevant Crashes

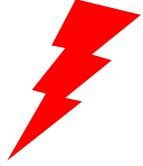
Pros:

- Very fast
- Same process space easy for instrumentation/code coverage
- Easy to know which input caused crash/replicate

Cons:

- Different from end exploit
- Might be naïve about initialization routines





The Fuzzing Cycle

We have an attack vector – but what should we send?

Identify an
attack vector

Generate a
Corpus of
Inputs

Create a
Fuzzing
Harness

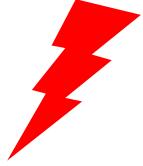
Fuzz and
Produce
Crashes

Identify
Relevant
Crashes

Sending totally random data is not likely to produce meaningful crashes

- Exception handlers
- Input validation

We need to identify examples of valid mach messages (e.g. “corpus building”)



The Fuzzing Cycle

Prep-Work

Identify an attack vector

Generate a Corpus of Inputs


Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

A number of things to take into consideration when we start debugging on MacOS

1. Setting up a MacOS virtual machine
2. Disabling System Integrity Protection (SIP)
 - csrutil disable
3. Disabling ReportCrash
4. Disabling Sleep
 - systemsetup -setsleep Never
5. Much more information provided: [Jeremy Brown - Summer of Fuzz: MacOS - DEF CON 29 AppSec Village](#)



The Fuzzing Cycle

Finding the Mach Message Handler

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

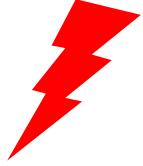
1

Find a mach service of interest

In our case, will be services sandboxed processes can communicate with

Let's focus on **com.apple.audio.coreaudiod**

- Handles all interactions with audio hardware
- Privileged process
- Allowed to send mach messages from many processes



The Fuzzing Cycle

Finding the Mach Message Handler

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes


2

Find the binary that implements the mach service

- **com.apple.audio.coreaudiod** registered with **launchd**
- Spawns **/usr/sbin/coreaudiod**
- Mach server handled by CoreAudio Framework

```
(lldb) image list
[ 0] D5BCB621-948E-308C-AF2C-88489D5569FA 0x000000010f332000 /usr/sbin/coreaudiod
[ 1] BB7A0970-8C62-3DCE-A7A2-5CEC9C501F11 0x00007ff80894f000 /usr/lib/dyld
[ 2] 66BBA3CA-BCE1-32F8-8269-99FAC92469FC 0x00007ff8123d6000 /System/Library/PrivateFrameworks/caulk.framework/Versions/A/caulk
[ 3] 97A3CD09-7112-376C-9613-7F38D4CF8C41 0x00007ff80ac99000 /System/Library/Frameworks/CoreAudio.framework/Versions/A/CoreAudio
[ 4] BEB5FC0B-7196-3C1D-A59A-F62ADA98F592 0x00007ff808ce4000 /System/Library/Frameworks/CoreFoundation.framework/Versions/A/CoreFoundation
```

```
fuzzychicken@Fuzzys-Mac HALB_MIGServer_server % stat /System/Library/Frameworks/CoreAudio.framework/Versions/A/CoreAudio
stat: cannot stat '/System/Library/Frameworks/CoreAudio.framework/Versions/A/CoreAudio': No such file or directory
```



The Fuzzing Cycle

Finding the Mach Message Handler

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

3

Extract the binary from the dyld shared cache

- **Dyld shared cache:** Starting with Big Sur, most framework binaries are not on disk
- We can extract them!
- <https://github.com/keith/dyld-shared-cache-extractor>

```
build> ./dyld-shared-cache-extractor /System/Volumes/Preboot/Cryptexes/OS/System/Library/dyld/dyld_shared_cache_x86_64h extracted-binaries
extracted 0/2505
extracted 1/2505
extracted 2/2505
extracted 3/2505
extracted 4/2505
extracted 5/2505
extracted 6/2505
extracted 7/2505
extracted 9/2505
extracted 8/2505
```

dyld-shared-cache-extractor Public

keith Add new use to README and format commands b... c28b25a · last month 19 Commits

.clang-format Replace with C (#2) 3 years ago

.gitignore Ignore more build directories last month

CMakeLists.txt Update extractor to use builtin or cust... last month

LICENSE Add LICENSE and README 3 years ago

README.md Add new use to README and format c... last month

dyld-shared-cache-extractor.c Validate passed library path exists last month

README MIT license

dyld-shared-cache-extractor

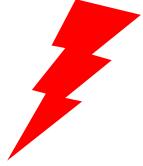
About A CLI for extracting libraries from Apple's dyld shared cache file

Readme MIT license Activity 353 stars 7 watching 30 forks Report repository

Releases 3

Support system dsc_extractor... on Dec 11, 2023 + 2 releases

Packages



The Fuzzing Cycle

Finding the Mach Message Handler

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness


Fuzz and Produce Crashes


Identify Relevant Crashes

4

Find function implementing mach receive functionality

- Wait, isn't this just **mach_msg()**?
 - Non-blocking, traps to kernel when a message is received
- Need to perform kernel debugging if we want to intercept incoming mach messages
 - This has been done:
<https://www.fortinet.com/blog/threat-research/inspect-mach-messages-in-macos-kernel-mode--part-ii--sniffing-th>
- Kernel debugging cons:
 - We see all mach messages, difficult to isolate target process
 - Two-machine debugging required
- Is there an easier way?





The Fuzzing Cycle

Finding the Mach Message Handler

Identify an attack vector

Generate a Corpus of Inputs


Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Mach Interface Generator (MIG)

- Apple provides MIG to more easily write RPC handlers and clients
- Interface Definition Language (IDL) compiler
- Abstracts much of the mach IPC layer away
- What if we searched for MIG-generated routines and dumped their incoming mach messages?





The Fuzzing Cycle

Finding the Mach Message Handler

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

3

Find function implementing mach receive functionality

- Hopper script:

<https://github.com/knightsc/hopper/blob/master/scripts/MIG%20Detect.py>

```
build/framework-binaries> nm -m ./System/Library/Frameworks/CoreAudio.framework/Versions/A/CoreAudio | grep -i subsystem
```

```
                (undefined) external _CACentralStateDumpRegisterSubsystem (from AudioToolboxCore)
```

```
00007ff8401adec0 (__DATA_CONST,__const) non-external _HALC_HALB_MIGClient_subsystem  
00007ff8401adf00 (__DATA_CONST,__const) non-external _HALS_HALB_MIGServer_subsystem
```





The Fuzzing Cycle

Finding the Mach Message Handler

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

3 Find function implementing mach receive functionality

_HALS_HALB_MIGServer_subsystem

- Function lookup table

```
; Attributes: bp-based frame  
; _int64 __fastcall HALB_MIGServer_server(mach_msg_header_t * __thiscall _HALB_MIGServer_server proc near  
push    rbp  
mov     rbp, rsp  
mov     eax, [rdi]  
and     eax, 1Fh  
mov     [rsi], eax  
mov     eax, [rdi+8]  
mov     [rsi+8], eax  
mov     dword ptr [rsi+4], 24h ; '$'  
xor     eax, eax  
mov     [rsi+0Ch], eax  
mov     ecx, [rdi+mach_msg_header_t.msgh_id]  
add     ecx, 64h ; 'd'  
mov     [rsi+14h], ecx  
mov     [rsi+10h], eax  
mov     ecx, -1010000  
add     ecx, [rdi+mach_msg_header_t.msgh_id] ; Get the msg ID Get msg ID  
cmp     ecx, 3Dh ; '='  
ja     short loc_7FF81DB61D64
```

Incoming msg (rdi)

```
mov    ecx, ecx  
lea    rcx, [rcx+rcx*4]  
lea    rdx, _HALS_HALB_MIGServer_subsystem  
mov    rcx, [rdx+rcx*8+28h] ; Index into function handler based on msg ID  
test   rcx, rcx  
jz     short loc_7FF81DB61D64
```

Get subsystem offset

Call function

```
call   rcx      ; Call the function  
mov    eax, 1  
jmp    short loc_7FF81DB61D79
```

```
loc_7FF81DB61D64:  
mov    rcx, cs:7FF85D276FF8h  
mov    rcx, [rcx]  
mov    [rsi+18h], rcx  
mov    dword ptr [rsi+20h], 0FFFFFED1h
```

```
loc_7FF81DB61D79:  
pop    rbp  
ret  
_HALB_MIGServer_server endp
```



The Fuzzing Cycle

Finding the Mach Message Handler

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

3 Find function implementing mach receive functionality

_HALS_HALB_MIGServer_subsystem

- Function lookup table

Function name

f	XObject_PropertyListener
f	XIOContext_PauseIO
f	XIOContext_ResumeIO
f	XIOContext_StopIO
f	XObject_GroupPropertyListener
f	XObject_GroupPropertyListener_Sync
f	XSystem_Open
f	XSystem_Close
f	XSystem_GetOpenInode
f	XSystem_CreateIOContext
f	XSystem_DestroyIOContext
f	XSystem_CreateMetaDevice
f	XSystem_DestroyMetaDevice
f	XSystem_ReadSetting
f	XSystem_WriteSetting
f	XSystem_DeleteSetting
f	XIOContext_SetClientControlPort
f	XIOContext_Start
f	XIOContext_Stop
f	XObject_HasProperty
f	XObject_IsPropertySettable
f	XObjectGetPropertyData
f	XObjectGetPropertyData_DI32
f	XObjectGetPropertyData_DI32_QI32
f	XObjectGetPropertyData_DAI32
f	XObjectGetPropertyData_DAI32_QAI32
f	XObjectGetPropertyData_DCFString
f	XObjectGetPropertyData_DCFString_QI32
f	XObjectGetPropertyData_DF32
f	XObjectGetPropertyData_DF32_QF32
f	XObjectGetPropertyData_DF64
f	XObjectGetPropertyData_DAF64
f	XObjectGetPropertyData_DPLList
f	XObjectGetPropertyData_DCFURL
f	XObject_SetPropertyData
f	XObject_SetPropertyData_DI32
f	XObject_SetPropertyData_DF32


RPC Functions

```
; Attributes: bp-based frame

_XSystem_Open proc near

var_D0= qword ptr -0D0h
var_C0= byte ptr -0C0h
var_B8= byte ptr -0B8h
var_B0= byte ptr -0B0h
var_A0= audit_token_t ptr -0A0h
var_80= qword ptr -80h
var_78= qword ptr -78h
var_70= xmmword ptr -70h
var_60= xmmword ptr -60h
buf= byte ptr -50h
var_30= qword ptr -30h

push    rbp
mov     rbp, rsp
push    r15
push    r14
push    r13
push    r12
push    rbx
sub    rsp, 0A8h
mov     r12, rsi
mov     rax, cs:7FF85D277498h
mov     rax, [rax]
mov     [rbp+var_30], rax
mov     ebx, 0FFFFFED0h
cmp    dword ptr [rdi], 0
jns    loc_7FF81DB4A118
```



The Fuzzing Cycle

Generate a Corpus of Inputs

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

I wrote a simple script to hook onto the message handler using LLDB

```
fuzzychicken@Fuzzys-Mac mach-fuzzing % sudo python3 subsystem_mach_msg_dumper.py -h
INFO Adding the LLDB Python library to PATH...
usage: subsystem_mach_msg_dumper.py [-h] -p PID -m MODULE -f FUNCTION
```

Attach to a process and dump a mach message passed to a specified function. The mach message should be passed as the first argument.

options:

```
-h, --help            show this help message and exit
-p PID, --pid PID    Process ID to attach to.
-m MODULE, --module MODULE
                      Module loaded by the process.
-f FUNCTION, --function FUNCTION
                      Function to set a breakpoint on.
```



The Fuzzing Cycle

Generate a Corpus of Inputs

Identify an attack vector


Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes





The Fuzzing Cycle

What is a Fuzzing Harness?

Identify an attack vector

Generate a Corpus of Inputs


Create a Fuzzing Harness


Fuzz and Produce Crashes

Identify Relevant Crashes



A **fuzzing harness** is code that allows you to send input through an attack vector.
(Call a desired function)





The Fuzzing Cycle

Calling the Target Function

Identify an attack vector

Generate a Corpus of Inputs


Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Target Function: `_HALB_MIGServer_server`

- Simple on Windows:
 - `HMODULE hModule = LoadLibrary("libexample.dll")`
 - `pFunction = GetProcAddress(hModule, "DesiredFunction")`
- On MacOS, similar:
 - `void *lib_handle = dlopen("libexample.dylib", RTLD_LAZY)`
 - `pFunction = dlsym(lib_handle, "DesiredFunction")`
- What if the symbol isn't exported?
- Write your own Mach-O symbol parser
 - A talk for another time ☺



The Fuzzing Cycle

Calling the Target Function

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Target

- Simulate the target function:
 - `HALB_MIGServer_server`
- On Mac OS X:
 - `msg_bits`: 4370
 - `msg_size`: 48
 - `msg_remote_port`: 106187
 - `msg_local_port`: 67075
 - `msg_voucher_port`: 0
 - `msg_id`: 1010013
- On Linux:
 - `msg_bits`: 18
 - `msg_size`: 36
 - `msg_remote_port`: 106187
 - `msg_local_port`: 0
 - `msg_voucher_port`: 0
 - `msg_id`: 1010113

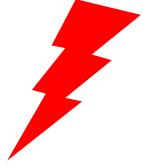
```
fuzzychicken@Fuzzys-Mac Release % ./Harness -l /System/Library/Frameworks/CoreAudio.framework/Versions/A/CoreAudio -s _HALB_MIGServer_server -f ~/mach-fuzzing/subsystem_messages/CoreAudio/HALB_MIGServer_server/a37747c4812a6baf1e4f5e723d78d4c3
```

Mach Message (Input)

Function (Attack Vector)

Return Mach Message

Calendar



The Fuzzing Cycle

What is a Fuzzer?

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness


Fuzz and Produce Crashes

Identify Relevant Crashes



A **fuzzer** is a program that generates inputs to be sent to a system and monitors for crashes.





The Fuzzing Cycle

What is a Fuzzer?


Identify an attack vector


Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes





The Fuzzing Cycle

The Need For Code Coverage

Identify an attack vector


Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

```
1 void process_string(const char *input_string) {
2     if (strlen(input_string) > 3) {
3         if (strlen(input_string == 6)) {
4             if (input_string[0] == 's') {
5                 if (strstr(input_string, "secret") != NULL) {
6                     int *ptr = NULL;
7                     *ptr = 1; // CRASH
8                 }
9             }
10            }
11        }
12    }
```



The Fuzzing Cycle

What is Code Coverage


Identify an attack vector

Generate a Corpus of Inputs


Create a Fuzzing Harness

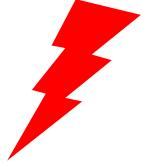
Fuzz and Produce Crashes

Identify Relevant Crashes



Code coverage traces a program's execution flow to identify new code paths.





The Fuzzing Cycle

How Do We Determine Code Coverage?


Identify an attack vector

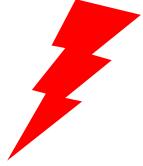
Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes







The Fuzzing Cycle

How Do We Determine Code Coverage?

Use instrumentation to monitor basic block execution

- Simple with source code:
 - AFL++ (<https://github.com/AFLplusplus/AFLplusplus>)
 - LibFuzzer (<https://llvm.org/docs/LibFuzzer.html>)
 - gCov (<https://gcc.gnu.org/onlinedocs/gcc/Gcov.html>)
- More difficult with black box binaries:
 - Frida (<https://frida.re/>)
 - TinyInst (<https://github.com/googleprojectzero/TinyInst>)
- Interpreting code coverage:
 - LightHouse for IdaPro/BinaryNinja (<https://github.com/gaasedelen/lighthouse>)





The Fuzzing Cycle

Actually Fuzzing!

Identify an attack vector

Generate a Corpus of Inputs


Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

My fuzzing setup

- Jackalope Fuzzer
(<https://github.com/googleprojectzero/Jackalope>)
- Enable Apple's GuardMalloc
 - Restricted pages placed surrounding all allocations
 - `DYLD_INSERT_LIBRARIES=/usr/lib/libgmalloc.dylib`
- TinyInst for dynamic instrumentation to dump coverage
- LightHouse to interpret code coverage



The Fuzzing Cycle

Actually Fuzzing!

Identify an attack vector


Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

```
fuzzychicken@Fuzzys-Mac Release % ./fuzzer -in ../../../../modified_msg_ids  
-out audio-startup-modified-ids -t 200 -t1 5000 -delivery file -instrument_module CoreAudio -target_module Harness -target_method _fuzz -nargs 1 -iterations 1000 -persist -loop -cmp_coverage -generate_unwind -dump_coverage -target_env DYLD_INSERT_LIBRARIES=/usr/lib/libmalloc.dylib --  
./Harness -f @@ -l /System/Library/Frameworks/CoreAudio.framework/Versions/A/CoreAudio -s _HALB_MIGServer_server  
Fuzzer version 1.00  
63 input files read  
Running input sample ../../../../../../modified_msg_ids/1010000  
GuardMalloc[Harness-3598]: Allocations will be placed on 16 byte boundaries.  
GuardMalloc[Harness-3598]: - Some buffer overruns may not be noticed.  
GuardMalloc[Harness-3598]: - Applications using vector instructions (e.g., SSE) should work.  
GuardMalloc[Harness-3598]: version 064555.99.1  
Instrumented module CoreAudio, code size: 7462910 GuardMalloc  
  
Total execs: 2  
Unique samples: 0 (0 discarded)  
Crashes: 0 (0 unique)  
Hangs: 0  
Offsets: 0  
Execs/s: 2  
GuardMalloc[Harness-3599]: Allocations will be placed on 16 byte boundaries.  
GuardMalloc[Harness-3599]: - Some buffer overruns may not be noticed.  
GuardMalloc[Harness-3599]: - Applications using vector instructions (e.g., SSE) should work.  
GuardMalloc[Harness-3599]: version 064555.99.1  
Instrumented module CoreAudio, code size: 7462910  
Exception at address 0x7ff85d79c63b C++ Exception  
Access address: 0x108d8000  
Exception in instrumented module CoreAudio 0x7ff81bbcd000  
Code before:  
47 ff ff c6 05 7e 72 bd 01 01  
Code after:  
41 89 5c 24 20 48 8b 05 b1 29 11 fe 48 8b 00 49  
GuardMalloc[Harness-3600]: Allocations will be placed on 16 byte boundaries.
```



The Fuzzing Cycle

Identify an
attack vector

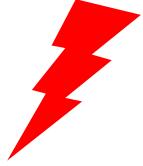
Generate a
Corpus of
Inputs

Create a
Fuzzing
Harness

Fuzz and
Produce
Crashes

Identify
Relevant
Crashes

Improving Our Fuzzing



The Fuzzing Cycle

Regularly Check Code Coverage

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

We can learn a lot from the code paths our fuzzer does and doesn't take

Goal: Cover as much of the binary as possible!

```
loc_7FF81DB4D953:
mov    r15d, [rdi+20h]
movups xmm0, xmmword ptr [rax+14h]
movups xmm1, xmmword ptr [rax+24h]
movaps [rbp+var_60], xmm1
movaps [rbp+var_70], xmm0
mov    rax, [rdi+24h]
mov    ecx, [rdi+2Ch]
movaps xmmword ptr [rbp+var_B0.val], xmm0
movaps xmmword ptr [rbp+var_B0.val+10h], xmm1
mov    [rbp+var_90], rax
mov    [rbp+var_88], ecx
xorps xmm0, xmm0
lea    rdx, [rbp+var_C0]
movaps xmmword ptr [rdx], xmm0
lea    rdi, [rbp+var_80]
xor    esi, esi
call   HALS_System::GetInstance(HALS_System::GetInstanceSetting, std::shared_ptr<HALS_System>)
mov    rdi, [rbp+var_80] ; this
test   rdi, rdi
jz    loc_7FF81DB4DA59
```

```
loc_7FF81DB4DA59: ; oslog
mov    rdi, cs:7FF85D277768h
mov    esi, 10h ; type
call   _os_log_type_enabled
test   al, al
jz    short loc_7FF81DB4DAB8
```


We always go this way!

```
lea    rsi, [rbp+var_B0] ; unsigned int
call   HALS_System::CopyClientByAuditToken(audit_token_t
mov    r14, rax
test   rax, rax
jz    loc_7FF81DB4DAEE
```

```
lea    r8, [rbp+buf] ; buf
mov    dword ptr [r8], 8200202h
lea    rax, aHalsMigserverC ; "HALS_MIGServer.cpp"
mov    [r8+4], rax
mov    word ptr [r8+0Ch], 400h
mov    dword ptr [r8+0Eh], 5F4h
lea    rdi, dso ; dso
mov    rsi, cs:7FF85D277768h ; log
lea    rcx, a25s5dHalsObjec ; "%25s:-%5d HALS_Object_HasProperty: t
mov    edx, 10h ; type
mov    r9d, 12h ; size
call   __os_log_impl
```


```
loc_7FF81DB4DAEE: ; oslog
mov    rdi, cs:7FF85D277768h
mov    esi, 10h ; type
call   _os_log_type_enabled
test   al, al
jz    short loc_7FF81DB4DB4D
```


Error: There is no system



The Fuzzing Cycle

Initialization: It Can Be a Big Problem








The Fuzzing Cycle

Initialization: It Can Be a Big Problem

- Identify an attack vector
- Generate a Corpus of Inputs
- Create a Fuzzing Harness
- Fuzz and Produce Crashes
- Identify Relevant Crashes





The Fuzzing Cycle

Initialization: It Can Be a Big Problem


I noticed my fuzzer was not getting deep into the code...

Why? The `coreaudiod` process performs initialization and setup for the **CoreAudio** Framework

It normally calls all these functions!

We want our fuzzer to be within the same process space as the initialized state

Coverage Overview			
Cov %	Func Name	Address	Blocks Hit
100.00	<code>_CFURLCreateWi...</code>	0x7FF81DB620...	1 / 1
100.00	<code>_CFURLGetString</code>	0x7FF81DB620...	1 / 1
100.00	<code>_CFURLGetTypeID</code>	0x7FF81DB620...	1 / 1
100.00	<code>_CFUUIDGetCons...</code>	0x7FF81DB620...	1 / 1
100.00	<code>_CFUUIDGetUUID...</code>	0x7FF81DB620...	1 / 1
100.00	<code>_IOConnectMapM...</code>	0x7FF81DB621...	1 / 1
100.00	<code>_IOConnectSetN...</code>	0x7FF81DB621...	1 / 1
100.00	<code>_IOIteratorNext</code>	0x7FF81DB621...	1 / 1
100.00	<code>_IOMainPort</code>	0x7FF81DB621...	1 / 1
100.00	<code>_IONotificatio...</code>	0x7FF81DB621...	1 / 1
100.00	<code>_IONotificatio...</code>	0x7FF81DB621...	1 / 1
100.00	<code>_IOObjectConfo...</code>	0x7FF81DB621...	1 / 1
100.00	<code>_IOObjectRelease</code>	0x7FF81DB621...	1 / 1
100.00	<code>_IOObjectRetain</code>	0x7FF81DB621...	1 / 1
100.00	<code>_IOPM GetUserAc...</code>	0x7FF81DB621...	1 / 1
100.00	<code>_IOPMScheduleU...</code>	0x7FF81DB622...	1 / 1
100.00	<code>_IORRegisterFor...</code>	0x7FF81DB622...	1 / 1
100.00	<code>_IORRegistryEnt...</code>	0x7FF81DB622...	1 / 1
100.00	<code>_IORRegistryEnt...</code>	0x7FF81DB622...	1 / 1
100.00	<code>_IORRegistryEnt...</code>	0x7FF81DB622...	1 / 1
100.00	<code>_IORRegistryEnt...</code>	0x7FF81DB622...	1 / 1
100.00	<code>_IORRegistryGet...</code>	0x7FF81DB622...	1 / 1
100.00	<code>_IOServiceAddI...</code>	0x7FF81DB622...	1 / 1
100.00	<code>_IOServiceAddM...</code>	0x7FF81DB622...	1 / 1
100.00	<code>_IOServiceMatc...</code>	0x7FF81DB622...	1 / 1



The Fuzzing Cycle

Initialization: It Can Be a Big Problem

Identify an attack vector

Generate a Corpus of Inputs


Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Possible Approaches:

1. Determine the initialization functions that need to be called, then call them within your fuzzing harness before starting fuzzing
 - Great when the initialization is simple
2. Modify a process you know performs proper initialization so it spawns your fuzzer
 - Hook a function through binary patching, `LD_PRELOAD`, detours, etc.
 - Instrument the binary using Frida or TinyInst



The Fuzzing Cycle

Initialization: It Can Be a Big Problem

Identify an attack vector

Generate a Corpus of Inputs


Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Possible Approaches:

1. Determine the initialization functions that need to be called, then call them within your fuzzing harness before starting fuzzing
 - Great when the initialization is simple
2. Modify a process you know performs proper initialization so it spawns your fuzzer
 - Hook a function through binary patching, `LD_PRELOAD`, detours, etc.
 - **Instrument the binary using Frida or TinyInst**



The Fuzzing Cycle

Identifying Where to Instrument

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness


Fuzz and Produce Crashes


Identify Relevant Crashes

coreaudiod loops continuously after performing the proper setup

`_CFRunLoopRunInMode` function

What if we changed that function so it calls our target?





The Fuzzing Cycle

Injecting Our Fuzzing Harness with Frida

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes


Identify Relevant Crashes

```
'use strict';

const harnessLib = Module.load('libharness.dylib');
const fuzz = new NativeFunction(harnessLib.getExportByName('fuzz'), 'void', []);

Interceptor.attach(Module.findExportByName('CoreFoundation', 'CFRunLoopRunInMode'), {
    onEnter: function(args) {
        console.log('CFRunLoopRunInMode called. Starting Fuzzer!');
        fuzz();
    },
    onLeave: function(retval) {
        console.log('Returning');
    }
});

console.log('Hook applied successfully.');
```



The Fuzzing Cycle

Injecting Our Fuzzing Harness with Frida

Identify an attack vector


Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

```
(venv) → mach-fuzzing git:(main) ✘ frida -l frida-fuzzer-injector.js -p 7412
    /---| Frida 16.3.3 - A world-class dynamic instrumentation toolkit
    | (_| |
    > _ | Commands:
    /_/_ | help      -> Displays the help system
    . . . . object?   -> Display information about 'object'
    . . . . exit/quit -> Exit
    . . . .
    . . . . More info at https://frida.re/docs/home/
    . . . .
    . . . . Connected to Local System (id=local)
Attaching...
Hook applied successfully.
[Local::PID::7412 ]-> CFRunLoopRunInMode called. Starting Fuzzer!
Process terminated
[Local::PID::7412 ]->
Thank you for using Frida!
(venv) → mach-fuzzing git:(main) ✘
```



The Fuzzing Cycle

Injecting Our Fuzzing Harness with Frida

Identify an attack vector


Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

```
(lldb) b fuzz
Breakpoint 3: where = libfuzzer.dylib`fuzz, address = 0x00000001000c0830
(lldb) c
Process 8157 resuming
(venv) → mach-fuzzing git:(n) Process 8157 stopped
    |---|
    |   |   Frida 16.3.3 - A
    |   |   frame #0: 0x00000001000c0830 libfuzzer.dylib`fuzz
    |   |   libfuzzer.dylib`fuzz:
    |   |   Commands:
    |   |       help      -> 0x1000c0830 <+0>: push    rbp
    |   |       object?   -> 0x1000c0831 <+1>: mov     rbp, rsp
    |   |       exit/quit -> 0x1000c0834 <+4>: sub     rsp, 0x30
    |   |       exit/quit -> 0x1000c0838 <+8>: mov     qword ptr [rbp - 0x8], rdi
    |   |       [lldb] x $rdi
    |   |       More info at ht0x103636580: 12 00 00 00 80 00 00 00 8f 8e 02 00 03 dc 00 00 ..... .
    |   |       0x103636590: 00 00 13 c2 87 69 0f 00 33 00 00 00 00 00 00 00 .....i..3.....
    |   |       Connected to Local System (id=local)
Attaching...
Hook applied successfully.
[Local::PID::7412 ]-> CFRunLoopRunInMode called. Starting Fuzzer!
Process terminated
[Local::PID::7412 ]->
Thank you for using Frida!
(venv) → mach-fuzzing git:(main) ×
```



The Fuzzing Cycle

Injecting a Full Fuzzer with Frida (In Progress)

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes


Identify Relevant Crashes


We want to inject a full fuzzer:

- Corpus management
- Mutation/input generation
- Code coverage
- Crash identification

This is not a trivial task

- LibAFL is a library that allows for custom fuzzer design using state-of-the art approaches
 - <https://github.com/AFLplusplus/LibAFL>
- Can help us run fuzzers in interesting places





The Fuzzing Cycle

Exploitable Versus Non-Exploitable Crashes

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes


Identify Relevant Crashes

Exploitable:

- Crash on write
- Crash on execution
- Illegal instruction
- Heap corruption abort
- Stack trace contains **free**, **malloc**, etc.

Likely Non-Exploitable:

- Crash on read (could be used to leak memory, though)
- Handled exception
- Null pointer dereferences
- Stack recursion



The Fuzzing Cycle

Exploitable Versus Non-Exploitable Crashes

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Exploitable:


- Crash on write
- Crash on execution
- Illegal instruction
- Heap corruption abort
- Stack trace contains **free**, **malloc**, etc.

Likely Non-Exploitable:

- Crash on read (could be used to leak memory, though)
- Handled exception
- Null pointer dereferences
- Stack recursion

Useful Tools:

- Apple's CrashWrangler (<https://developer.apple.com/library/archive/technotes/tn2334/index.html>)
- CrashMon (<https://github.com/ant4g0nist/crashmon>)



The Fuzzing Cycle

Exploitable Versus Non-Exploitable Crashes

Identify an attack vector

Generate a Corpus of Inputs

Create a Fuzzing Harness

Fuzz and Produce Crashes

Identify Relevant Crashes

Exploitable:

- Crash on write
- Crash on execution
- Illegal instruction
- Heap corruption abort
- Stack trace contains **free**, **malloc**, etc.

Likely Non-Exploitable:

- Crash on read (could be used to leak memory, though)
- Handled exception
- Null pointer dereferences
- Stack recursion

Useful Tools:

- Apple's CrashWrangler (<https://developer.apple.com/library/archive/technotes/tn2334/index.html>)
- CrashMon (<https://github.com/ant4g0nist/crashmon>)

Crash Reproducibility

- Should be able to run input through harness and reproduce the crash



What We've Covered

- A crash course on fuzzing and Mach IPC mechanisms
- A walkthrough of the attack process:
 - Identifying an attack vector
 - Generating a corpus of fuzzing inputs
 - Writing a custom fuzzing harness
 - Fuzzing and producing crashes
 - Crash triaging
- Common pitfalls and things to consider
- Inspired you to do vulnerability research!



Personal Advice

- Try to get your fuzzer as close to the target function as possible
- The harder something is to fuzz, the more likely there will be vulnerabilities
- Let your code coverage inform your fuzzing modifications



FUZZING TAKEAWAYS

Questions

Thank You!

Twitter: @dillon_franke

Blog: <https://dillonfrankesecurity.com>